

StyCLIP: StyleTransfer in Few-Shot Learning

Emanuele Fontana

Università degli Studi di Bari Aldo Moro

`e.fontana7@studenti.uniba.it`

Introduction & Motivation

- **Context:** Rapid digitization of art collections demands automated artist classification systems for large-scale analysis.
- **Challenge:** Few-shot visual artist classification – identifying artist by their style with limited labeled examples per artist.
- **Problem:** Current image classifiers are based on semantic features, not style (e.g., CLIP).
- **Hypothesis:** Explicitly integrating style representations (Gram matrices) into CLIP could enhance artist-specific pattern recognition.
- **Contribution:** StyCLIP – a CLIP extension incorporating dedicated style-aware features for improved few-shot artist classification.

Related Work

Style Transfer

- Gatys et al. [3]: Deep networks encode style in intermediate layers; Gram matrices capture style.

Few-Shot Learning (FSL) & Meta-Learning

- Prototypical Networks [6]: Learn to learn from small samples.

Contrastive Language-Image Pretraining (CLIP)

- CLIP [5]: Shared embedding space for images and text.
- TIP-Adapter [8]: Non-parametric cache model for FSL.
- APE [9]: Adaptive prior refinement of CLIP features.
- GDA-CLIP [7]: Gradient-based adaptation.
- TIMO [4]: Mutual guidance between text/image modalities.

Background: TIMO (Text-Image Mutual Guidance Optimization)




Figure: TIMO Architecture Overview

Background: TIMO (Text-Image Mutual Guidance Optimization)

Two Main Branches

Text-Guided Image (TGI): Uses textual representations to refine and guide visual features

Image-Guided Text (IGT): Finds optimal weights for textual prompts using support visual features, dynamically adapting textual descriptions to image

ArtGraph Dataset Separation





Figure: ArtGraph Datasets Overview

Episode Example for FSL



Figure: Example episode: 5-Way, 4-Shots ,10 query examples

StyCLIP: Architecture Overview



StyCLIP: Key Components

- **Core Idea:** Extend CLIP with explicit style-aware features via Gram matrices
- **Three Key Components:**
 1. **Gram Matrix Extraction:** Capture style correlations from intermediate visual layers (ResNet-50 blocks)
 2. **Style Feature Projection:** Reduce dimensionality of Gram matrix representations
 3. **Fusion Adapter:** Integrate semantic (CLIP) and style features through bottleneck architecture

StyCLIP: Training Strategy

- **Frozen Backbone:** Original CLIP (RN50) visual and text encoders are frozen.
- **Trainable Components:** Only the new style projection layers and fusion adapter are optimized.
- **Rationale:**
 - Preserve pre-trained semantic knowledge from CLIP.
 - Reduce computational cost.
 - Enable efficient fine-tuning on artistic datasets.

Hyperparameter Tuning

- Optuna [1] used for:
 - Fusion bottleneck dimension {64, 128, 256}
 - Dropout rate [0.0, 0.5]
 - Learning rate log-uniform [10^{-5} , 10^{-3}]
 - Weight decay log-uniform [10^{-6} , 10^{-3}]

Experimental Setup

Meta-Learning Training for StyCLIP

- Prototypical network approach.
- Episodic training: 10-way ($N=10$ artists), K -shot ($K=1, 4$) classification, 3-query ($Q=3$)

Evaluation Protocol

- **Method:** Using every CLIP-based method (e.g., TIMO, TIP-Adapter) with every backbone.
- N -way classification (all available artists in the test subset).
- Tested several times across $K \in \{1, 2, 4, 8, 16\}$ shots for support set.
- **Metrics:** Accuracy, Macro-Precision, Macro-Recall, Macro-F1.

Key Evaluation Results

Main Finding: Integrating explicit style-aware features (StyCLIP) *consistently underperforms* standard CLIP variants.

- Standard CLIP with RN50 backbone is the best performing backbone.
- StyCLIP backbones (StyCLIP_FSL1, StyCLIP_FSL4) significantly degrade performance for all FSL methods.
- TIMO and TIMO-S are the best FSL methods when using the standard RN50 backbone.

Example: 16-Shot Classification Accuracy (%)

Model	Backbone	ACC (%)
TIMO	StyCLIP_FSL1	37.00
TIMO	StyCLIP_FSL4	43.27
TIMO	RN50	73.36
TIMO_S	StyCLIP_FSL1	39.20
TIMO_S	StyCLIP_FSL4	48.75
TIMO_S	RN50	73.46

Statistical Analysis Summary

- T-tests confirmed the significance of performance differences.
- **Key Observation:** Standard CLIP (RN50) significantly outperforms StyCLIP-trained backbones across nearly all configurations and few-shot methods.
- Example: For 16-shot with TIMO_S:
 - RN50 (73.46%) vs StyCLIP_FSL4 (48.75%) $\rightarrow \approx 25$ percentage point difference, statistically significant.

$\Rightarrow$ The architectural changes in StyCLIP for explicit style modeling did not yield benefits and often hindered performance.

Conclusion & Discussion

- **Main Finding:** Integrating explicit style-aware features via Gram matrices (StyCLIP) did **not improve** few-shot artist classification performance.
- **Performance:** Standard CLIP (RN50 backbone) consistently outperformed StyCLIP variants.
- **Implication:** Explicit Gram matrix-based style modeling might interfere with CLIP's pre-trained semantic representations rather than enhance them for this task.
- **Future Directions:** Alternative approaches like attention-based style integration or domain-specific fine-tuning (without architectural changes) might be more fruitful.

Acknowledgments

- This work was proposed by PhD student **Nicola Fanelli** and **Prof. Giovanna Castellano** at the University of Bari.

Code Availability

The code for this project is available on GitHub:

`https://github.com/Fonty02/ComputerVision/tree/main/TIM0`

Thank You! 

Appendix: Hyperparameter Optimization Highlights

StyCLIP FSL1 (1-shot trained)

- Bottleneck: 64
- LR: 6.01×10^{-5}
- Dropout: 0.0
- Weight Decay: 4.83×10^{-4}
- Valid Acc: 57.07%

StyCLIP FSL4 (4-shot trained)

- Bottleneck: 128
- LR: 1.02×10^{-5}
- Dropout: 0.1
- Weight Decay: 1.03×10^{-6}
- Valid Acc: 60.53%

References I



Takuya Akiba, Shotaro Sano, Toshihiko Yanase, Takeru Ohta, and Masanori Koyama.

Optuna: A next-generation hyperparameter optimization framework.

In *Proceedings of the 25th ACM SIGKDD international conference on knowledge discovery & data mining*, pages 2623–2631, 2019.



Giovanna Castellano, Vincenzo Digeno, Giovanni Sansaro, and Gennaro Vessio.

Artgraph, March 2022.



Leon A. Gatys, Alexander S. Ecker, and Matthias Bethge.

Image style transfer using convolutional neural networks.

In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2016.



Yayuan Li, Jintao Guo, Lei Qi, Wenbin Li, and Yinghuan Shi.

Text and image are mutually beneficial: Enhancing training-free few-shot classification with clip.

In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 5039–5047, 2025.

References II



Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al.

Learning transferable visual models from natural language supervision.

In *International conference on machine learning*, pages 8748–8763. PmLR, 2021.



Jake Snell, Kevin Swersky, and Richard Zemel.

Prototypical networks for few-shot learning.

In I. Guyon, U. Von Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 30. Curran Associates, Inc., 2017.



Zhengbo Wang, Jian Liang, Lijun Sheng, Ran He, Zilei Wang, and Tieniu Tan.

A hard-to-beat baseline for training-free clip-based adaptation.

arXiv preprint arXiv:2402.04087, 2024.

References III



Renrui Zhang, Wei Zhang, Rongyao Fang, Peng Gao, Kunchang Li, Jifeng Dai, Yu Qiao, and Hongsheng Li.

Tip-adapter: Training-free adaption of clip for few-shot classification.

In *European conference on computer vision*, pages 493–510. Springer, 2022.



Xiangyang Zhu, Renrui Zhang, Bowei He, Aojun Zhou, Dong Wang, Bin Zhao, and Peng Gao.

Not all features matter: Enhancing few-shot clip with adaptive prior refinement.

In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 2605–2615, 2023.

The End